

ECE 595: Homework 4

Rahul Deshmukh, Class ID: 58

(Spring 2019)

In collaboration with: Murali Krishnan & Anmol Singh Guram

Exercise 1

a Convergence of Logistic Regression:

(i) Given: The two datasets are linearly separable.

To Prove: The magnitude of slope and intercept parameters \mathbf{w} and w_0 would tend to ∞

Proof: If the data is linearly separable then we would want the Loss function to be zero. The Loss function for Logistic regression is given by:

$$\begin{aligned} J(\boldsymbol{\theta}) &= - \sum_j (y_j \log(h_{\boldsymbol{\theta}}(\mathbf{x}_j)) + (1 - y_j) \log(1 - h_{\boldsymbol{\theta}}(\mathbf{x}_j))) \\ &= - \sum_{\mathbf{x}_j \in C_1} \log(h_{\boldsymbol{\theta}}(\mathbf{x}_j)) - \sum_{\mathbf{x}_j \in C_0} \log(1 - h_{\boldsymbol{\theta}}(\mathbf{x}_j)) \\ \text{where } h_{\boldsymbol{\theta}}(\mathbf{x}_j) &= \frac{1}{1 + e^{\boldsymbol{\theta}^T \mathbf{x}_j}} \end{aligned}$$

We can observe that:

$$\begin{aligned} h_{\boldsymbol{\theta}}(\mathbf{x}_j) &< 1 \text{ for finite value of } \boldsymbol{\theta} \\ \Rightarrow \log(h_{\boldsymbol{\theta}}(\mathbf{x}_j)) &< 0 \text{ \& } \log(1 - h_{\boldsymbol{\theta}}(\mathbf{x}_j)) < 0 \\ \Rightarrow J(\boldsymbol{\theta}) &> 0 \end{aligned}$$

Also, we can say the following:

$$\begin{aligned} \mathbf{x}_j \in C_1 \quad \lim_{\boldsymbol{\theta} \rightarrow \infty} h_{\boldsymbol{\theta}}(\mathbf{x}_j) &= 1 \\ \mathbf{x}_j \in C_0 \quad \lim_{\boldsymbol{\theta} \rightarrow \infty} h_{\boldsymbol{\theta}}(\mathbf{x}_j) &= 0 \\ \Rightarrow \lim_{\boldsymbol{\theta} \rightarrow \infty} J(\boldsymbol{\theta}) &= 0 \end{aligned}$$

Therefore, we will attain the minimum zero loss only when $\boldsymbol{\theta} \rightarrow \infty$. This means that for any solution $\boldsymbol{\theta}_k$ we can obtain a better solution $\alpha \boldsymbol{\theta}_k$ where $\alpha > 0$ with a smaller loss value. This indicates that the logistic regression will suffer from nonconvergence in case of linearly separable data.

To Prove: Gradient descent iterates would not converge in a finite number of steps if we have the stopping criteria as $\|\boldsymbol{\theta}^{(k+1)} - \boldsymbol{\theta}^{(k)}\| = 0$

Proof:

$$\begin{aligned}\boldsymbol{\theta}^{(k+1)} &= \boldsymbol{\theta}^{(k)} - \alpha_k \left(\sum_{n=1}^N (h_{\boldsymbol{\theta}^{(k)}}(\mathbf{x}_n) - y_n) \mathbf{x}_n \right) \\ \boldsymbol{\theta}^{(k+1)} - \boldsymbol{\theta}^{(k)} &= -\alpha_k \left(\sum_{n=1}^N (h_{\boldsymbol{\theta}^{(k)}}(\mathbf{x}_n) - y_n) \mathbf{x}_n \right) \\ \text{if } \|\boldsymbol{\theta}^{(k+1)} - \boldsymbol{\theta}^{(k)}\| &= 0 \Rightarrow h_{\boldsymbol{\theta}^{(k)}}(\mathbf{x}_n) = y_n \\ &\Rightarrow \text{only when } \boldsymbol{\theta} \rightarrow \infty \\ &\Rightarrow k \rightarrow \infty \\ &\text{ie never converges}\end{aligned}$$

(ii) When we have $\|\mathbf{w}\| < c_1$ and $|w_0| < c_2$ for some $c_1, c_2 > 0$, then we can say that $\boldsymbol{\theta}$ is no longer unbounded and cannot reach to ∞ . Then these limits will be used as a termination criteria and thus we will have a converged solution.

Some other ways to counter the non-convergence issue can be to add a regularization term ($\lambda \|\boldsymbol{\theta}\|^2$) to the loss function and solve the unconstrained problem. **Or** To Rather minimize the function $\|\boldsymbol{\theta}\|^2$ subject to the constraint $J(\boldsymbol{\theta}) < \epsilon$.

(iii) We don't face the issue of non-convergence of other linear classifiers when the data is linearly separable this is because for the other methods we can attain a zero loss for linearly separable data within finite number of iterations (which we have proved in class for several methods)

b Proof of convergence of online mode of perceptron algorithm:

Given: $\mathbf{w}^{(k+1)} = \mathbf{w}^{(k)} + \alpha_k y_j \mathbf{x}_j$ with $\alpha_k = 1$ for $j \in \mathcal{M}_k$

(i) To prove: The move from $\mathbf{w}^{(k)}$ to $\mathbf{w}^{(k+1)}$ is in the right direction i.e. decreases the misclassification error.

We need to show that $y_j(\mathbf{w}^{(k+1)})^T \mathbf{x}_j > y_j(\mathbf{w}^{(k)})^T \mathbf{x}_j$

Proof:

$$\begin{aligned}
 y_j(\mathbf{w}^{(k+1)})^T \mathbf{x}_j &= y_j(\mathbf{w}^{(k)} + \alpha_k y_j \mathbf{x}_j)^T \mathbf{x}_j \\
 &= y_j(\mathbf{w}^{(k)})^T \mathbf{x}_j + y_j(\alpha_k y_j \mathbf{x}_j)^T \mathbf{x}_j \\
 &= y_j(\mathbf{w}^{(k)})^T \mathbf{x}_j + (\alpha_k y_j^2 \|\mathbf{x}_j\|^2) \\
 &\geq y_j(\mathbf{w}^{(k)})^T \mathbf{x}_j \dots\dots\dots \text{as } \alpha_k y_j^2 \|\mathbf{x}_j\|^2 \geq 0 \\
 \Rightarrow y_j(\mathbf{w}^{(k+1)})^T \mathbf{x}_j &\geq y_j(\mathbf{w}^{(k)})^T \mathbf{x}_j \\
 &\text{hence proved}
 \end{aligned}$$

(ii) Proof of convergence:

Given: $\rho = \min_j y_j(\mathbf{w}^*)^T \mathbf{x}_j$

(1) To prove: $(\mathbf{w}^{(k)})^T \mathbf{w}^* \geq k\rho$ (Please note: this is a tighter bound than the homework prompt)

Proof:

$$\begin{aligned}
 (\mathbf{w}^{(k)})^T \mathbf{w}^* &= (\mathbf{w}^{(k-1)} + y_j \mathbf{x}_j)^T \mathbf{w}^* \\
 &= (\mathbf{w}^{(k-1)})^T \mathbf{w}^* + (y_j \mathbf{x}_j)^T \mathbf{w}^* \\
 &= (\mathbf{w}^{(k-1)})^T \mathbf{w}^* + y_j (\mathbf{w}^*)^T \mathbf{x}_j \\
 &\text{by definition: } y_j (\mathbf{w}^*)^T \mathbf{x}_j \geq \rho \\
 &\geq (\mathbf{w}^{(k-1)})^T \mathbf{w}^* + \rho \\
 &\text{By induction we can say:} \\
 (\mathbf{w}^{(k)})^T \mathbf{w}^* &\geq (\mathbf{w}^{(0)})^T \mathbf{w}^* + k\rho \\
 &\text{where } \mathbf{w}^{(0)} = \mathbf{0} \\
 \Rightarrow (\mathbf{w}^{(k)})^T \mathbf{w}^* &\geq k\rho \\
 &\text{hence proved}
 \end{aligned}$$

(2) Given: $R = \max_j \|\mathbf{x}_j\|_2$, To prove: $\|\mathbf{w}^{(k)}\|_2^2 \leq k^2 R^2$

Proof:

$$\|\mathbf{w}^{(k)}\| = \|\mathbf{w}^{(k-1)} + y_j \mathbf{x}_j\|$$

using Triangle inequality we get:

$$\leq \|\mathbf{w}^{(k-1)}\| + \|y_j \mathbf{x}_j\|$$

$$\leq \|\mathbf{w}^{(k-1)}\| + \|\mathbf{x}_j\| \dots\dots\dots \text{as } y_j = \pm 1$$

By induction we have:

$$\leq \|\mathbf{w}^{(0)}\| + k\|\mathbf{x}_j\|$$

where $\mathbf{w}^{(0)} = \mathbf{0}$

$$\leq k\|\mathbf{x}_j\|$$

$$\leq kR$$

$$\Rightarrow \|\mathbf{w}^{(k)}\|^2 \leq k^2 R^2$$

hence proved

(3) To Prove: $\frac{(\mathbf{w}^{(k)})^T \mathbf{w}^*}{\|\mathbf{w}^{(k)}\|_2} \geq \frac{\rho}{R}$ and $k \leq \frac{R^2 \|\mathbf{w}^*\|}{\rho^2}$ (Please note: this is a tighter bound than the homework prompt)

$$(\mathbf{w}^{(k)})^T \mathbf{w}^* \geq k\rho \text{ (proved in part (1))}$$

$$\|\mathbf{w}^{(k)}\| \leq kR \text{ (proved in part (2))}$$

$$\Rightarrow \frac{(\mathbf{w}^{(k)})^T \mathbf{w}^*}{\|\mathbf{w}^{(k)}\|_2} \geq \frac{k\rho}{kR} = \frac{\rho}{R}$$

hence proved

Now let us try to prove that the sequence $\mathbf{w}^{(k)}$ converges to \mathbf{w}^* . Therefore we need to prove the following:

$$\exists \epsilon > 0 \mid \|\mathbf{w}^{(k)} - \mathbf{w}^*\| < \epsilon$$

Proof:

$$\begin{aligned}
\|\mathbf{w}^{(k)} - \mathbf{w}^*\|^2 &= \|\mathbf{w}^{(k-1)} + \alpha_k y_j \mathbf{x}_j - \mathbf{w}^*\|^2 \\
&= \|\mathbf{w}^{(k-1)} - \mathbf{w}^* + \alpha_k y_j \mathbf{x}_j\|^2 \\
&= \|\mathbf{w}^{(k-1)} - \mathbf{w}^*\|^2 + \|\alpha_k y_j \mathbf{x}_j\|^2 + 2\alpha_k (\mathbf{w}^{(k-1)} - \mathbf{w}^*)^T y_j \mathbf{x}_j
\end{aligned}$$

we know that at k-1 iteration the misclassification error is greater than 0 ie $(\mathbf{w}^{(k-1)})^T y_j \mathbf{x}_j \leq 0$

$$\begin{aligned}
&\leq \|\mathbf{w}^{(k-1)} - \mathbf{w}^*\|^2 + \alpha_k^2 \|y_j \mathbf{x}_j\|^2 - 2\alpha_k (\mathbf{w}^*)^T y_j \mathbf{x}_j \\
&\leq \|\mathbf{w}^{(k-1)} - \mathbf{w}^*\|^2 + \alpha_k (\alpha_k \|y_j \mathbf{x}_j\|^2 - 2(\mathbf{w}^*)^T y_j \mathbf{x}_j) \\
&\text{if } \alpha_k \|y_j \mathbf{x}_j\|^2 - 2(\mathbf{w}^*)^T y_j \mathbf{x}_j \leq 0
\end{aligned}$$

$$\text{ie } \alpha_k \leq \frac{2(\mathbf{w}^*)^T y_j \mathbf{x}_j}{\|y_j \mathbf{x}_j\|^2}$$

$$\text{choosing } \alpha_k = \frac{(\mathbf{w}^*)^T y_j \mathbf{x}_j}{\|y_j \mathbf{x}_j\|^2} \text{ we get:}$$

$$\|\mathbf{w}^{(k)} - \mathbf{w}^*\|^2 \leq \|\mathbf{w}^{(k-1)} - \mathbf{w}^*\|^2 - \frac{((\mathbf{w}^*)^T y_j \mathbf{x}_j)^2}{\|y_j \mathbf{x}_j\|^2}$$

By induction we get:

$$\|\mathbf{w}^{(k)} - \mathbf{w}^*\|^2 \leq \|\mathbf{w}^{(0)} - \mathbf{w}^*\|^2 - k \frac{((\mathbf{w}^*)^T y_j \mathbf{x}_j)^2}{\|y_j \mathbf{x}_j\|^2}$$

where $\mathbf{w}^{(0)} = \mathbf{0}$

$$\|\mathbf{w}^{(k)} - \mathbf{w}^*\|^2 \leq \|\mathbf{w}^*\|^2 - k \frac{((\mathbf{w}^*)^T y_j \mathbf{x}_j)^2}{\|y_j \mathbf{x}_j\|^2} = \epsilon^2$$

Therefore, The sequence the sequence $\mathbf{w}^{(k)}$ converges to \mathbf{w}^* if $\epsilon^2 > 0$:

$$\Rightarrow \|\mathbf{w}^*\|^2 - k \frac{((\mathbf{w}^*)^T y_j \mathbf{x}_j)^2}{\|y_j \mathbf{x}_j\|^2} > 0$$

$$\|\mathbf{w}^*\|^2 > k \frac{((\mathbf{w}^*)^T y_j \mathbf{x}_j)^2}{\|y_j \mathbf{x}_j\|^2}$$

$$k < \frac{\|\mathbf{w}^*\|^2 \|y_j \mathbf{x}_j\|^2}{((\mathbf{w}^*)^T y_j \mathbf{x}_j)^2}$$

$$k < \frac{\|\mathbf{w}^*\|^2 \|\mathbf{x}_j\|^2}{((\mathbf{w}^*)^T y_j \mathbf{x}_j)^2}$$

using earlier definitions $\|\mathbf{x}_j\| \leq R$ and $(\mathbf{w}^*)^T y_j \mathbf{x}_j \geq \rho$

$$k \leq \frac{\|\mathbf{w}^*\|^2 R^2}{\rho^2}$$

hence proved

c Soft-Margin SVM:

$$\begin{aligned} & \underset{\boldsymbol{\theta}, \boldsymbol{\xi}}{\operatorname{argmin}} \frac{1}{2} \|\mathbf{w}\|_2^2 + \frac{C}{2} \|\boldsymbol{\xi}\|_2^2 \\ & \text{subject to } y_j g_{\boldsymbol{\theta}}(\mathbf{x}_j) \geq 1 - \xi_j \\ & \quad \xi_j \geq 0, j = 1, \dots, N \end{aligned}$$

We need to find the convex dual problem for it

(i) The constraint $\xi_j \geq 0$ can be removed without affecting the solution to the optimization problem, why? lets write ou the lagrangian of the function:

$$L(\theta, \lambda, \xi, \mu) = \frac{1}{2} \|\mathbf{w}\|_2^2 + \frac{C}{2} \|\boldsymbol{\xi}\|_2^2 + \sum_j (\lambda_j (1 - \xi_j - y_j (\mathbf{w}^T \mathbf{x}_j + w_0)) + \mu_j (-\xi_j))$$

Form complementary slackness condition we have:

$$\begin{aligned} \mu_j (\xi_j) &= 0 \\ \Rightarrow \mu_j &= 0, \xi_j \geq 0 \\ \text{or } \mu_j &\geq 0, \xi_j = 0 \end{aligned}$$

Therefore the constraint $\xi_j \geq 0$ is always satisfied and using the above fact we can simplify the lagrangian to:

$$L(\theta, \lambda, \xi) = \frac{1}{2} \|\mathbf{w}\|_2^2 + \frac{C}{2} \|\boldsymbol{\xi}\|_2^2 + \sum_j (\lambda_j (1 - \xi_j - y_j (\mathbf{w}^T \mathbf{x}_j + w_0)))$$

(ii) Finding the optimal solutions:

$$\begin{aligned} L(\theta, \lambda, \xi) &= \frac{1}{2} \|\mathbf{w}\|_2^2 + \frac{C}{2} \|\boldsymbol{\xi}\|_2^2 + \sum_j (\lambda_j (1 - \xi_j - y_j (\mathbf{w}^T \mathbf{x}_j + w_0))) \\ \nabla_{\mathbf{w}} L &= 0 \Rightarrow \mathbf{w}^* - \sum_j \lambda_j y_j \mathbf{x}_j = 0 \\ \Rightarrow \mathbf{w}^* &= \sum_j \lambda_j y_j \mathbf{x}_j \\ \nabla_{w_0} L &= 0 \Rightarrow \sum_j \lambda_j y_j = 0 \\ \nabla_{\boldsymbol{\xi}} L &= 0 \Rightarrow C \boldsymbol{\xi} - \boldsymbol{\lambda} = 0 \\ \Rightarrow C \boldsymbol{\xi} &= \boldsymbol{\lambda} \Rightarrow \boldsymbol{\xi}^* = \frac{\boldsymbol{\lambda}}{C} \end{aligned}$$

hence proved

(iii) Finding the convex dual of the problem:

$$\begin{aligned}
& \underset{\mathbf{w}, w_0}{\text{minimize}} \mathcal{L}(\mathbf{w}, w_0, \lambda) \\
&= \underset{\mathbf{w}, w_0}{\text{minimize}} \left\{ \frac{1}{2} \|\mathbf{w}\|_2^2 + \frac{C}{2} \|\boldsymbol{\xi}\|_2^2 + \max_{\lambda \geq 0} \left\{ \sum_j (\lambda_j (1 - \xi_j - y_j(\mathbf{w}^T \mathbf{x}_j + w_0))) \right\} \right\} \\
&= \underset{\mathbf{w}, w_0}{\text{minimize}} \{ \underset{\lambda}{\text{maximize}} \mathcal{L}(\mathbf{w}, w_0, \lambda) \} \\
&= \underset{\lambda}{\text{maximize}} \{ \underset{\mathbf{w}, w_0}{\text{minimize}} \mathcal{L}(\mathbf{w}, w_0, \lambda) \} \\
&= \underset{\lambda}{\text{maximize}} \{ \mathcal{L}(\mathbf{w}^*, w_0^*, \lambda) \} \\
&= \underset{\lambda}{\text{maximize}} \left\{ \frac{1}{2} \|\mathbf{w}^*\|_2^2 + \frac{C}{2} \|\boldsymbol{\xi}^*\|_2^2 + \sum_j (\lambda_j (1 - \xi_j^* - y_j(\mathbf{w}^{*T} \mathbf{x}_j + w_0^*))) \right\} \\
&= \underset{\lambda}{\text{maximize}} \left\{ + \frac{1}{2} \sum_i \sum_j \lambda_i \lambda_j y_i y_j \mathbf{x}_i^T \mathbf{x}_j + \frac{C}{2} \sum_k \frac{\lambda_k^2}{C} + \sum_j \lambda_j \right. \\
&\quad \left. - \sum_k \frac{\lambda_k^2}{C} - \sum_i \sum_j \lambda_i \lambda_j y_i y_j \mathbf{x}_i^T \mathbf{x}_j - w_0^* \sum_j (\lambda_j y_j) \right\} \\
&= \underset{\lambda}{\text{maximize}} \left\{ - \frac{1}{2} \sum_i \sum_j \lambda_i \lambda_j y_i y_j \mathbf{x}_i^T \mathbf{x}_j - \frac{1}{2} \sum_k \frac{\lambda_k^2}{C} + \sum_j \lambda_j \right\} \\
&\quad \text{subject to } \sum_j \lambda_j y_j = 0
\end{aligned}$$

hence proved

d Hard- Margin SVM:

$$\begin{aligned} \underset{\boldsymbol{\theta}}{\operatorname{argmin}} \frac{1}{2} \|\mathbf{w}\|_2^2 \\ \text{st } y_j g_{\boldsymbol{\theta}}(\mathbf{x}_j) \geq 0, j = 1, \dots, N \end{aligned}$$

(i) To prove: if the constraint is changed to $y_j g_{\boldsymbol{\theta}}(\mathbf{x}_j) \geq \gamma$ such that $\gamma \geq 0$ then the problem remains the same.

Proof:

$$\begin{aligned} \underset{\boldsymbol{\theta}}{\operatorname{argmin}} \frac{1}{2} \|\mathbf{w}\|_2^2 \\ y_j g_{\boldsymbol{\theta}}(\mathbf{x}_j) \geq \gamma \\ \Rightarrow y_j (\mathbf{w}^T \mathbf{x}_j + w_0) / \gamma \geq 1 \\ \Rightarrow y_j \left(\frac{\mathbf{w}^T}{\gamma} \mathbf{x}_j + \frac{w_0}{\gamma} \right) \geq 1 \end{aligned}$$

We can scale the objective function by a constant $\frac{1}{\gamma^2}$ without changing the problem

$$\begin{aligned} \Rightarrow \underset{\boldsymbol{\theta}}{\operatorname{argmin}} \frac{1}{2} \frac{\|\mathbf{w}\|_2^2}{\gamma^2} \\ y_j \left(\frac{\mathbf{w}^T}{\gamma} \mathbf{x}_j + \frac{w_0}{\gamma} \right) \geq 1 \\ \Rightarrow \underset{\hat{\boldsymbol{\theta}} = \frac{\boldsymbol{\theta}}{\gamma}}{\operatorname{argmin}} \frac{1}{2} \|\hat{\mathbf{w}}\|_2^2 \\ y_j (\hat{\mathbf{w}}^T \mathbf{x}_j + \hat{w}_0) \geq 1 \end{aligned}$$

This is the same problem as earlier. Hence proved.

(ii) To prove : SVM for two points $\mathbf{x}_1 \in C_+$ and $\mathbf{x}_2 \in C_-$ is solvable.

Proof:

The Dual Problem is:

$$\begin{aligned}
& \max_{\lambda \geq 0} \left\{ \sum_{j=1}^2 \lambda_j - \frac{1}{2} \sum_{i=1}^2 \sum_{j=1}^2 (\lambda_i \lambda_j y_i y_j \mathbf{x}_i^T \mathbf{x}_j) \right\} \\
& \text{st } \sum_{j=1}^2 \lambda_j y_j = 0 \\
& \Rightarrow \max_{\lambda \geq 0} \left\{ \sum_{j=1}^2 \lambda_j - \frac{1}{2} \sum_{i=1}^2 \sum_{j=1}^2 (\lambda_i \lambda_j y_i y_j \mathbf{x}_i^T \mathbf{x}_j) \right\} \\
& \text{st } \sum_{j=1}^2 \lambda_j y_j = 0 \\
& \Rightarrow \lambda_1 - \lambda_2 = 0 \\
& \Rightarrow \lambda_1 = \lambda_2 = \lambda \\
& \Rightarrow \max_{\lambda \geq 0} \left\{ 2\lambda - \frac{\lambda^2}{2} (\mathbf{x}_1^T \mathbf{x}_1 + \mathbf{x}_2^T \mathbf{x}_2 - 2\mathbf{x}_1^T \mathbf{x}_2) \right\} \\
& \text{taking derivative and equating to zero we get} \\
& 2 - \lambda(\|\mathbf{x}_1 - \mathbf{x}_2\|^2) = 0 \\
& \lambda = \frac{2}{\|\mathbf{x}_1 - \mathbf{x}_2\|^2} \geq 0 \\
& \mathbf{w}^* = \sum_{j=1}^2 \lambda_j y_j \mathbf{x}_j \\
& \Rightarrow \mathbf{w}^* = \lambda(\mathbf{x}_1 - \mathbf{x}_2) \\
& = \frac{2(\mathbf{x}_1 - \mathbf{x}_2)}{\|\mathbf{x}_1 - \mathbf{x}_2\|^2} \\
& w_0 = -\frac{\mathbf{w}^{*T}(\mathbf{x}_1 + \mathbf{x}_2)}{2} \\
& \Rightarrow w_0 = -\frac{1}{2} \lambda (\mathbf{x}_1 - \mathbf{x}_2)^T (\mathbf{x}_1 + \mathbf{x}_2) \\
& w_0 = -\frac{(\|\mathbf{x}_1\|^2 - \|\mathbf{x}_2\|^2)}{\|\mathbf{x}_1 - \mathbf{x}_2\|^2}
\end{aligned}$$

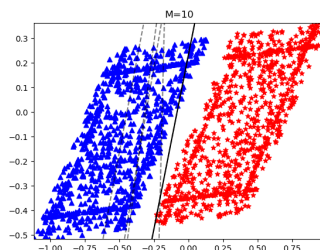
Hence proved.

Exercise 2

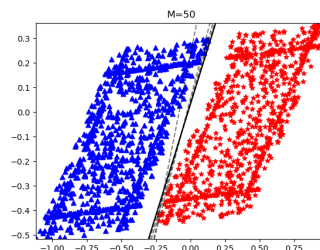
(a) Please find the code at page 15

(b) Please find the code at page 15

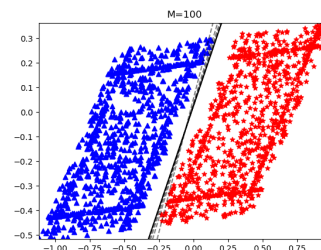
The Learning rate was chosen as 10^{-1}



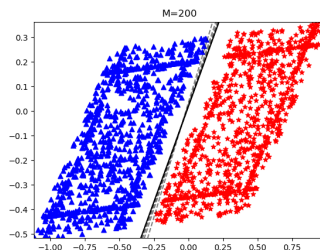
(a) M=10



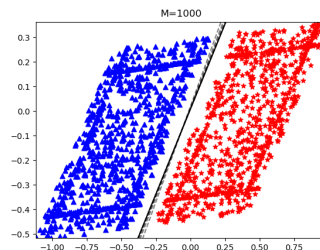
(b) M=50



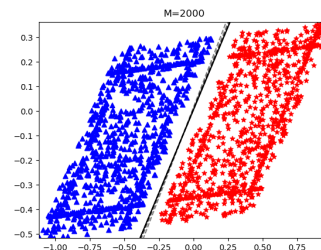
(c) M=100



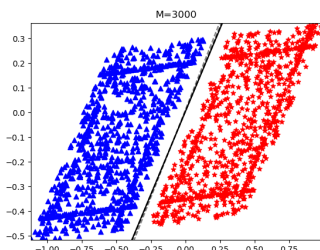
(d) M=200



(e) M=1000

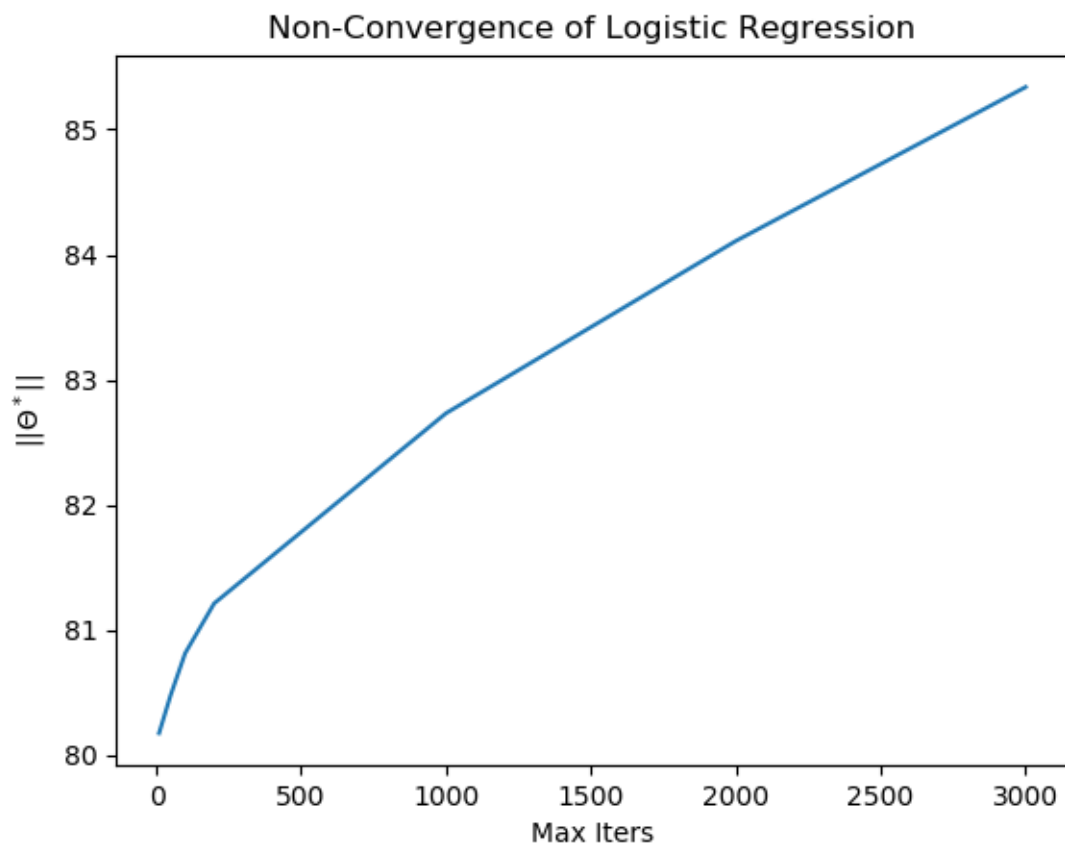


(f) M=2000



(g) M=3000

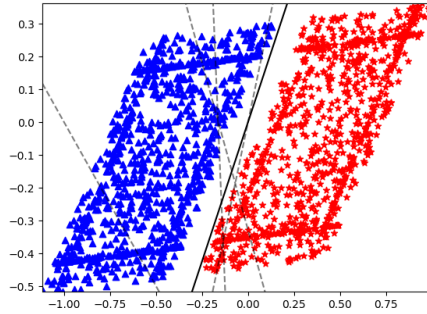
As can be observed, the decision boundary improves as the maximum number of iterations increases however, I observed that after attaining a good separating hyperplane there is no change in the decision boundary as we further increase the number of iterations.



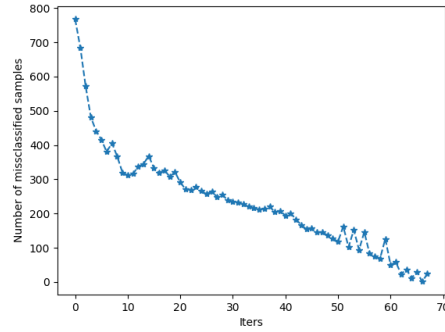
We can observe from the above plot that the magnitude of theta $\|\theta\|$ monotonically increases as the number of iterations increases this further confirms our proof in exercise 1(a).

b Perceptron: Please find the code at page 15

(i) Online Mode

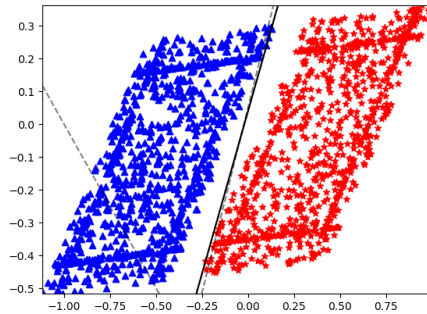


(h) decision boundary plot

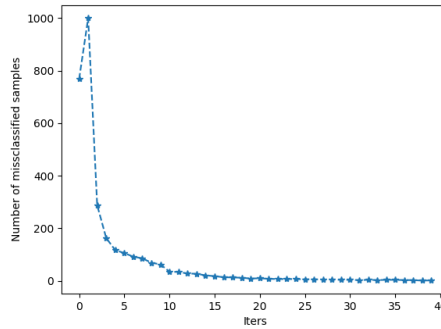


(i) Convergence

(i) Batch Mode



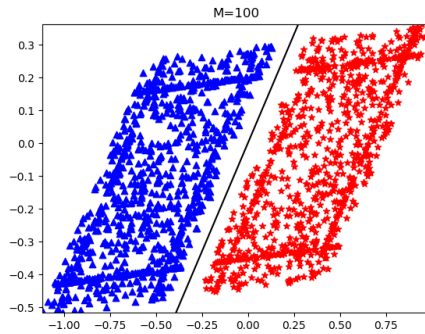
(j) decision boundary plot



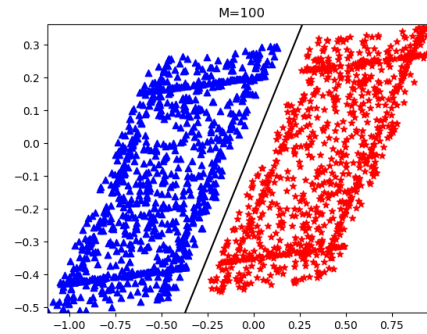
(k) Convergence

From the above plot for convergence of both the modes we observe that convergence is faster for Batch Mode. However, in terms of cost of computation online mode is cheaper than batch mode.

c SVM, Please find the code at page 15

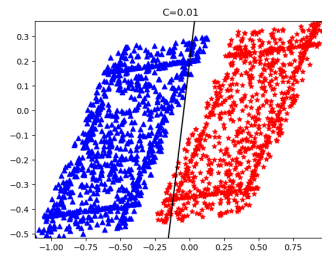


(l) Hard

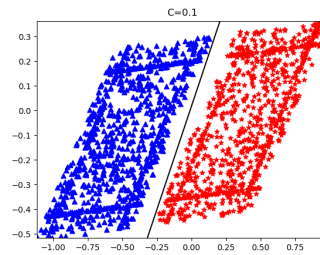


(m) soft

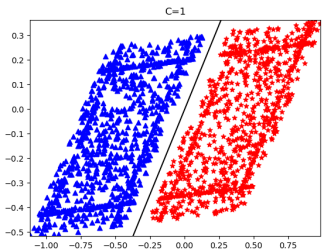
We can observe that the decision boundaries for the two cases are nearly. More insights can be gained from calculating the margin for the two cases. We can observe from the above plots that



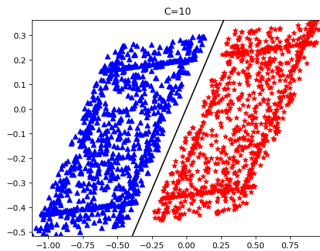
(n) $C=0.01$



(o) $C=0.1$



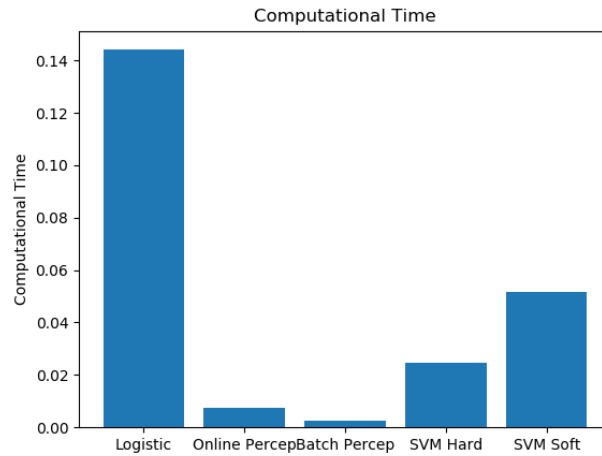
(p) $C=1$



(q) $C=10$

for the soft margin as C increases the decision boundary starts resembling the hard SVM decision boundary. Also, as the data is linearly separable with a good margin, therefore using a hard margin SVM is better suited.

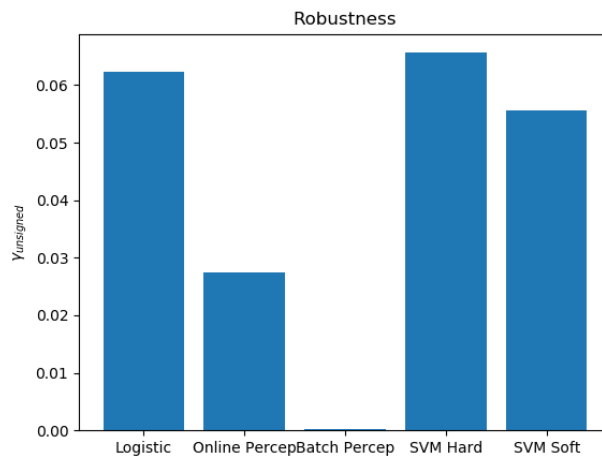
d Classification Error: For all of the classifiers we saw (from earlier plots) that we were able to get correct classification.



(r) Computational Time

Computational Time: We can observe that Logistic regression takes the maximum time to get a good solution as it requires more number of iterations. Also, The perceptron methods are cheaper than SVM because in perceptron we are not doing additional work of maximizing the margin.

Effect of Learning rate: It was observed that a smaller learning rate required more iterations to converge and a faster learning rate can give erroneous results.



(s) $\gamma_{unsigned}$

From the above plot we can observe that Perceptron Methods are the least robust and can be a poor method to classify a test dataset. Of all the classifiers SVM hard seems to be the most robust. This is because the dataset is linearly separable and has no outliers, had there been outliers then SVM soft would have been the best option.

Main Code for Exercise 2

Listing 1: Source Code

```
1  #!/usr/bin/env python3
2  # -*- coding: utf-8 -*-
3  """
4  Created on Mon Mar 18 16:47:50 2019
5  @author: rahul
6  """
7  # import libraries
8  import numpy as np , csv
9  from functions import *
10 import matplotlib.pyplot as plt
11 %% Read Data
12 #(a)
13 read_path = '../data/'
14 label_filename = 'hw04_labels.csv'
15 sample_filename = 'hw04_sample_vectors.csv'
16
17 samples = []
18 with open(read_path+sample_filename) as csvfile:
19     readcsv = csv.reader(csvfile , delimiter=',')
20     for row in readcsv:
21         vals = list(map(float , row))
22         samples.append( vals )
23 samples_copy = np.array( samples )
24 samples = np.ones((np.shape(samples_copy)[0] , np.shape(samples_copy)[1]+1))
25 samples[:, :-1] = samples_copy
26 samples = samples.T
27
28 labels = []
29 with open(read_path+label_filename) as csvfile:
30     readcsv = csv.reader(csvfile , delimiter=',')
31     for row in readcsv:
32         vals = float( row[0] )
33         labels.append( vals )
34 labels = np.array( labels )
35
36 scaled_labels = 2*labels-1 #percep labels in range -1 to +1
37 # Declare constants
38
39 rate = 0.1
40 %%
41 #(a) Logistic regression based classification
42 log_star_store = []
43 Ms = [10, 50, 100, 200, 1000, 2000, 3000]
44
45 for M in Ms:
46     [log_theta_star , log_theta_store] = logistic( samples , labels , rate , M )
47     log_star_store.append( np.linalg.norm( log_theta_star ) )
48     freq = int( 0.2*M )
49     # plotdata( samples , labels , log_theta_store , freq , M )
50     # plotdata_single( samples , labels , log_theta_star , M )
51     log_gamma = min( scaled_labels*(log_theta_star.T@samples)/np.linalg.norm(
52         log_theta_star ) )
53 #plt.figure()
54 #plt.plot( Ms , log_star_store )
```

```

55 #plt.xlabel('Max Iters')
56 #plt.ylabel('||$\Theta^*||$')
57 #plt.title('Non-Convergence of Logistic Regression')
58 #plt.show()
59
60 ##(b) Perceptron Online mode
61 M = 100 # max number of iters
62 freq = int(0.2*M)
63 scaled_labels = 2*labels-1 #percep labels in range -1 to +1
64 ##(i)
65 [ol_percp_theta_star, ol_percp_theta_store] = perceptron(samples, scaled_labels, rate, M,
    convplot=0)
66 #plotdata(samples, labels, ol_percp_theta_store, freq, [M])
67 ol_gamma = min((scaled_labels*(ol_percp_theta_star.T@samples))/np.linalg.norm(
    ol_percp_theta_star))
68
69
70 ##(ii)
71 [bt_percp_theta_star, bt_percp_theta_store] = perceptron(samples, scaled_labels, rate, M,
    convplot=0, online=False)
72 #plotdata(samples, labels, bt_percp_theta_store, freq)
73 bt_gamma = min(scaled_labels*(bt_percp_theta_star.T@samples)/np.linalg.norm(
    bt_percp_theta_star))
74
75 ##(c) SVM
76 #(i) Hard Margin
77 svm_hard_theta_star = SVM_hard(samples, scaled_labels)
78 #plotdata_single(samples, labels, svm_hard_theta_star, M)
79 hard_gamma = min((scaled_labels*(np.array(svm_hard_theta_star).T@samples))/np.linalg
    .norm(svm_hard_theta_star))
80
81 ##(ii) Soft Margin
82 C = 1
83 #for C in [10**(-2), 10**(-1), 1, 10]:
84 svm_soft_theta_star = SVM_soft_L1(samples, scaled_labels, C)
85 #plotdata_single(samples, labels, svm_soft_theta_star, C)
86 soft_gamma = min((scaled_labels*(np.array(svm_soft_theta_star).T@samples))/np.linalg
    .norm(svm_hard_theta_star))
87
88 #plt.figure()
89 #plt.bar(['Logistic', 'Online Percep', 'Batch Percep', 'SVM Hard', 'SVM Soft'], \
90 #        [log_gamma, ol_gamma, bt_gamma, hard_gamma, soft_gamma])
91 #plt.ylabel('$\gamma_{\text{unsigned}}$')
92 #plt.title('Robustness')
93 #plt.show()

```

Functions for Exercise 2

Listing 2: Functions

```

1 #!/usr/bin/env python3
2 #- coding: utf-8 -*-
3 """
4 Created on Mon Mar 18 17:07:13 2019
5
6 @author: rahul
7 """
8 # import libraries

```



```

9 import numpy as np
10 import matplotlib.pyplot as plt
11 import cvxpy as cvx
12 import time
13 """ SVM
14 # SVM Soft Margin
15 def SVM_soft_L1(samples, labels, C):
16     """
17     Input:
18     sample: np array in the format [[x1],[x2],...,[xn]] xi as col vectors
19     labels: np array [y1,y2,...,yn]
20     """
21     t_in = time.time()
22     dim, Nsamples = samples.shape
23     theta = cvx.Variable(dim) # declaring dimension of variable
24     lam = 1/C
25     Jump_term = cvx.max_elemwise(0, 1 - cvx.mul_elemwise(labels, (theta.T@samples).T))
26     Jump_term = cvx.sum_entries(Jump_term)
27     obj_expr = Jump_term + (lam/2)*(cvx.sum_squares(theta[:-1]))
28     obj = cvx.Minimize(obj_expr)
29     prob = cvx.Problem(obj)
30     prob.solve(solver = cvx.ECOS)
31     theta_star = theta.value
32     theta_star = theta_star.tolist()
33     theta_star = [x[0] for x in theta_star]
34     print('SVM soft time '+str(time.time() - t_in)+'\n')
35     return(theta_star)
36
37 # SVM Hard Margin
38 def SVM_hard(samples, labels):
39     """
40     Input:
41     sample: np array in the format [[x1],[x2],...,[xn]] xi as col vectors
42     labels: np array [y1,y2,...,yn]
43     """
44     t_in = time.time()
45     dim, Nsamples = samples.shape
46     theta = cvx.Variable(dim) # declaring dimension of variable
47     obj = cvx.Minimize(cvx.sum_squares(theta[:-1]))
48     temp = cvx.mul_elemwise(labels, (theta.T@samples).T)
49     const = [ temp >= 1 ]
50     prob = cvx.Problem(obj, const)
51     prob.solve(solver = cvx.ECOS)
52     theta_star = theta.value
53     #format change
54     theta_star = theta_star.tolist()
55     theta_star = [x[0] for x in theta_star]
56     print('SVM hard time '+str(time.time() - t_in)+'\n')
57     return(theta_star)
58
59
60 """ Perceptron Method
61 def percept_batch_tangent(theta, samples, true_labels):
62     """
63     Input:
64         theta: decision boundary
65         sample: np array in the format [[x1],[x2],...,[xn]] xi as col vectors
66         labels: np array [y1,y2,...,yn]
67     Out: J: Jacobian

```

```

68     """
69     # predict samples labels using theta
70     gx = theta.T@samples
71     ygx = np.multiply(true_labels, gx)
72     miss_idx = np.where(ygx<0)[0]
73     N_miss = miss_idx.size
74     all_miss_labels = true_labels[miss_idx]
75     all_miss_x = samples[:, miss_idx]
76     J = np.sum(all_miss_labels*all_miss_x, axis=1)
77     return(J, N_miss)
78
79 def percept_online_tangent(theta, samples, true_labels):
80     """
81     Input:
82         theta: decision boundary
83         sample: np array in the format [[x1],[x2],..., [xn]] xi as col vectors
84         labels: np array [y1,y2,..., yn]
85     Out: J: Jacobian
86     """
87     # predict samples labels using theta
88     gx = theta.T@samples
89     ygx = np.multiply(true_labels, gx)
90     miss_idx = np.where(ygx<0)[0]
91     N_miss = miss_idx.size
92     if N_miss>0:
93         picked_idx = miss_idx[np.random.permutation(N_miss)]
94         picked_idx = picked_idx[0]
95         J = true_labels[picked_idx]*samples[:, picked_idx]
96     else:
97         J = np.zeros_like(samples[:,0]) #just to pass some value
98     return(J, N_miss)
99
100 def perceptron(samples, labels, rate, Max_iter, convplot, online = True):
101     """
102     Input:
103         samples: np array in the format [[x1],[x2],..., [xn]] xi as col vectors
104         labels: np array [y1,y2,..., yn]
105         rate: learning rate or the step length
106         Max_iter: for the gradient descent
107         online: if online mode then True(default) else False
108     """
109     t_in = time.time()
110     xdim, Nsamples= samples.shape
111     theta_k = np.ones(xdim) # initial guess
112     theta_store = []
113     N_store = []
114     theta_store.append(np.copy(theta_k))
115     if online:
116         for k in range(Max_iter):
117             grad_k, N_miss = percept_online_tangent(theta_k, samples, labels)
118             if N_miss == 0:
119                 break
120             theta_k += rate*grad_k
121             theta_store.append(np.copy(theta_k))
122             N_store.append(N_miss)
123             print('online percep time '+str(time.time() - t_in)+'\n')
124     else:
125         for k in range(Max_iter):
126             grad_k, N_miss = percept_batch_tangent(theta_k, samples, labels)

```

```

127         if N_miss == 0:
128             break
129         theta_k += rate*grad_k
130         theta_store.append(np.copy(theta_k))
131         N_store.append(N_miss)
132         print('batch percep time '+str(time.time() - t_in)+'\n')
133     if convplot==1:
134         plt.plot(np.arange(len(N_store)),N_store, '*--')
135         plt.xlabel('Iters')
136         plt.ylabel('Number of missclassified samples')
137         plt.show()
138
139     return(theta_k, theta_store)
140
141 """ function for plotting data
142 def plotdata_single(samples, labels, theta_star, *argv):
143     """
144     Input:
145         sample: np array in the format [[x1],[x2],...,[xn]] xi as col vectors
146         labels: np array [y1,y2,...,yn]
147         theta_star: decision boundary params
148         N: frequency of plots for decision boundary
149     """
150
151     plt.figure()
152     #plot the training set
153     for i in range(np.shape(samples)[1]):
154         if labels[i] == 1:
155             plt.scatter(samples[0,i], samples[1,i], c='red', marker='*')
156         else:
157             plt.scatter(samples[0,i], samples[1,i], c='blue', marker='^')
158
159     # plot the decision boundary
160     x_min = min(samples[0,:])
161     x_max = max(samples[0,:])
162     y_min = min(samples[1,:])
163     y_max = max(samples[1,:])
164     x_line = np.linspace(x_min, x_max, 10)
165     y_line = -(theta_star[0]*x_line + theta_star[2])/theta_star[1]
166     plt.plot(x_line, y_line, 'k-')
167     plt.xlim(x_min, x_max)
168     plt.ylim(y_min, y_max)
169     if argv!=None:
170         plt.title('C='+str(argv[0]))
171     plt.show()
172     return()
173
174 def plotdata(samples, labels, theta_store, N, *argv):
175     """
176     Input:
177         sample: np array in the format [[x1],[x2],...,[xn]] xi as col vectors
178         labels: np array [y1,y2,...,yn]
179         theta_store: decision boundary params
180         N: frequency of plots for decision boundary
181     """
182
183     plt.figure()
184     #plot the training set
185     for i in range(np.shape(samples)[1]):

```

```

186         if labels[i] == 1:
187             plt.scatter(samples[0,i],samples[1,i],c='red',marker='*')
188         else:
189             plt.scatter(samples[0,i],samples[1,i],c='blue',marker='^')
190
191     # plot the decision boundary
192     x_min = min(samples[0,:])
193     x_max = max(samples[0,:])
194     y_min = min(samples[1,:])
195     y_max = max(samples[1,:])
196     x_line = np.linspace(x_min,x_max,10)
197     count=0
198     while count<len(theta_store):
199         theta_k = theta_store[count]
200         y_line = -(theta_k[0]*x_line +theta_k[2])/theta_k[1]
201         plt.plot(x_line,y_line,'k—',alpha=0.5)
202         count+=N
203     theta_k = theta_store[-1]
204     y_line = -(theta_k[0]*x_line +theta_k[2])/theta_k[1]
205     plt.plot(x_line,y_line,'k-')
206     plt.xlim(x_min,x_max)
207     plt.ylim(y_min,y_max)
208     if argv!=None:
209         plt.title('M='+str(argv[0]))
210     plt.show()
211     return()
212
213 %% Logistic regression function
214 # tangent for logistic function
215 def logistic_tangent(theta,samples,labels):
216     """
217     Input:
218         theta: decision boundary
219         sample: np array in the format [[x1],[x2],...,[xn]] xi as col vectors
220         labels: np array [y1,y2,...,yn]
221     Out: J: Jacobian
222     """
223     thetaTx = theta.T@samples
224     h_theta_x = 1/(1+np.exp(-thetaTx))
225     temp = (h_theta_x - labels)*samples
226     J = np.sum(temp,axis=1)
227     return(J)
228 # main logistic function
229 def logistic(samples,labels,rate,Max_iter):
230     """
231     Input:
232         samples: np array in the format [[x1],[x2],...,[xn]] xi as col vectors
233         labels: np array [y1,y2,...,yn]
234         rate: learninig rate or the step length
235         Max_iter: for the gradient descent
236     """
237     t_in = time.time()
238     xdim,Nsamples= samples.shape
239     theta_k = np.zeros(xdim) # initial guess
240     theta_store = []
241     theta_store.append(np.copy(theta_k))
242     for k in range(Max_iter):
243         theta_k -= rate*logistic_tangent(theta_k,samples,labels)
244         theta_store.append(np.copy(theta_k))

```

```
245     print( 'SVM logistic time ' + str( time.time() - t_in ) + '\n' )
246     return( theta_k , theta_store )
```